

Local-First Principles: a Behavioural Types Approach

Emilio Tuosto @ GSSI

joint work with

Roland Kuhn @ Actyx



and

Hernán Melgratti @ UBA



It-Matters
Lucca 11-12 July, 2023

– Prelude –

Take-away message

An approach to trade consistency for availability in systems of **asymmetric replicated peers**

Take-away message

An approach to trade consistency for availability in systems of **asymmetric replicated peers**

based on **local-first**'s principles to (re-)gain **consensus** ... eventually

Take-away message

An approach to trade consistency for availability in systems of **asymmetric replicated peers**

based on **local-first**'s principles to (re-)gain **consensus** ... eventually

formally supported by behavioural types

Take-away message

An approach to trade consistency for availability in systems of **asymmetric replicated peers**

based on **local-first**'s principles to (re-)gain **consensus** ... eventually

formally supported by behavioural types



- **swarm** = (**machines** + **local logs**) * **imaginary global log**
- **swarm protocols**: systems from an abstract **global** viewpoint
- enforce **good behaviour** via behavioural typing

Take-away message

An approach to trade consistency for availability in systems of **asymmetric replicated peers**

based on **local-first**'s principles to (re-)gain **consensus** ... eventually

formally supported by behavioural types




- **swarm** = (**machines** + **local logs**) * **imaginary global log**
- **swarm protocols**: systems from an abstract **global** viewpoint
- enforce **good behaviour** via behavioural typing



See our recent ECOOP 2023 paper
(to appear; extended version available at <https://arxiv.org/abs/2305.04848>)

Distributed coordination



An “old” problem

Distributed agreement

Distributed sharing

Security

Computer-assisted collaborative work

...

With some “solutions”


Centralisation points

Consensus protocols

Commutative replicated data types

...

Distributed coordination



Availability = Money

Kohavi et al. KDD'14

- Amazon sales down 1% if 100ms delay
- Google searches down 0.2% - 0.6% if 100-400ms delay
- Bing's revenue down ~1.5% if 250ms delay

An “old” problem

Distributed agreement
Distributed sharing
Security
Computer-assisted collaborative work
...

With some “solutions”

Centralisation points
Consensus protocols
Commutative replicated data types
...

A new (?) solution

What about using local-first principles?

Thou shall be autonomous

Thou shall collaborate

Thou shall recognise conflicts

Thou shall resolve conflicts

Thou shall be consistent

Plan of the talk

Some motivations

Our formalisation

Our typing discipline

Tool support

Open issues

– Motivations –

(the pictures are courtesy of Actyx AG)

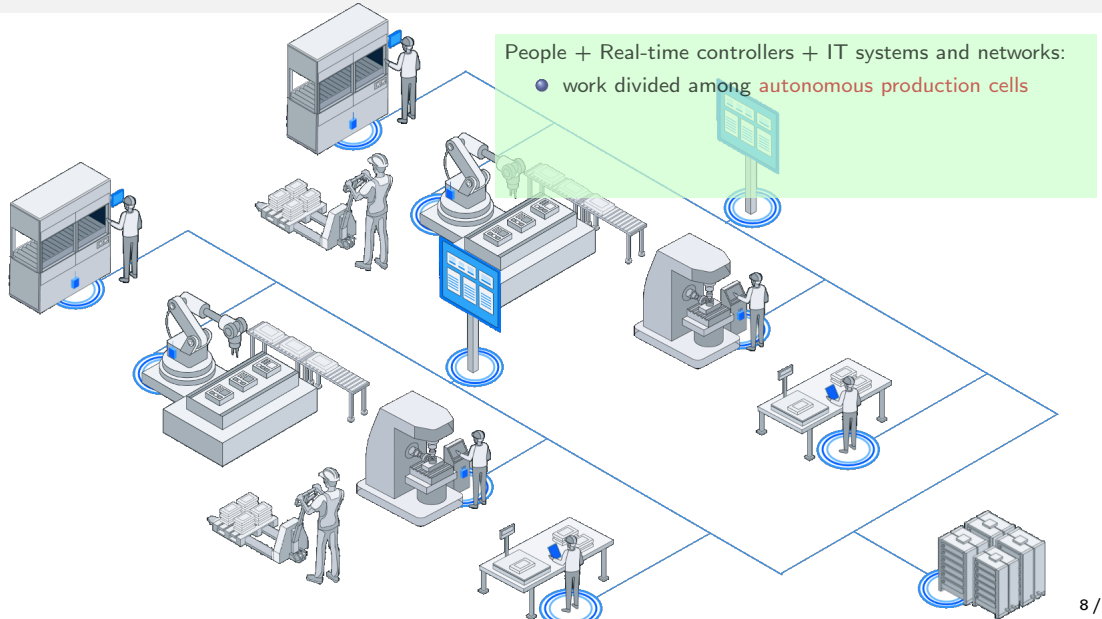


(the pictures are courtesy of Actyx AG)



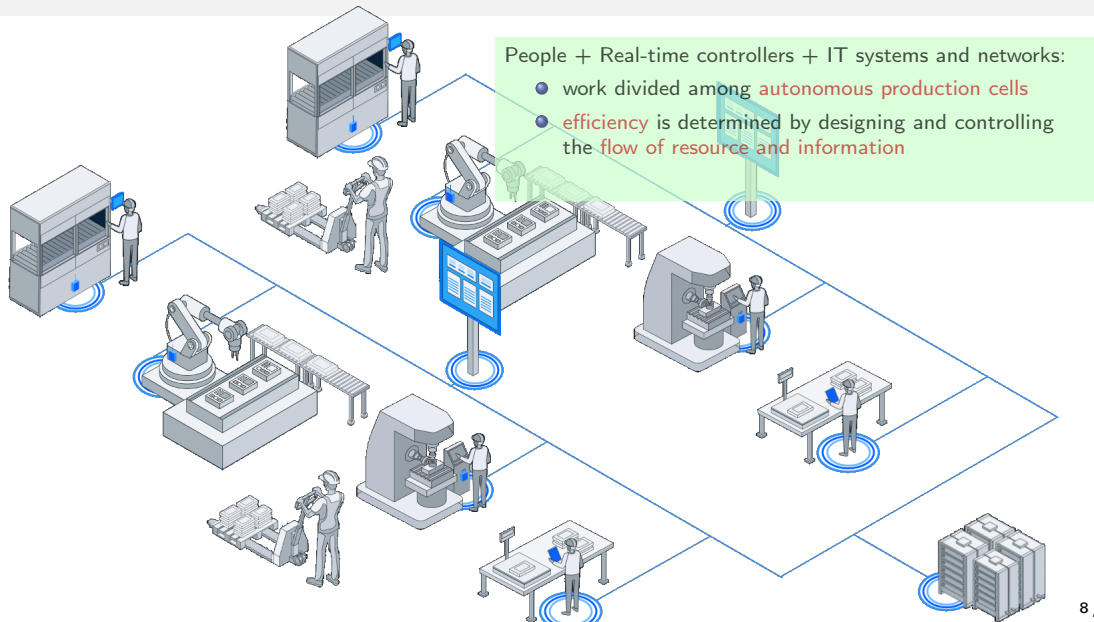
A collaborative environment and its execution model

(the pictures are courtesy of Actyx AG)



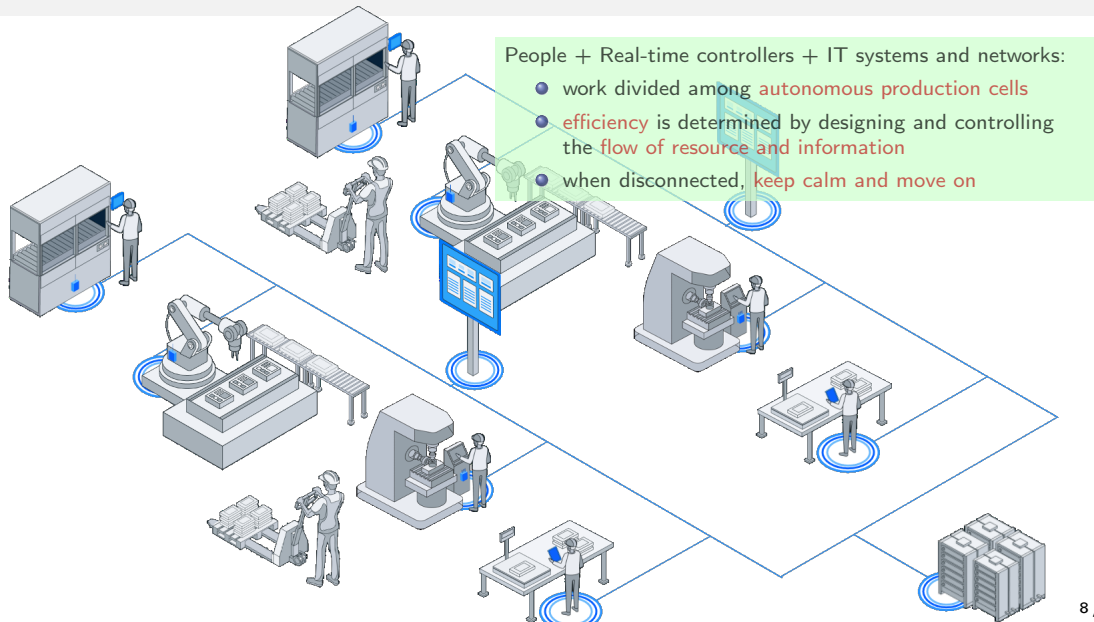
A collaborative environment and its execution model

(the pictures are courtesy of Actyx AG)



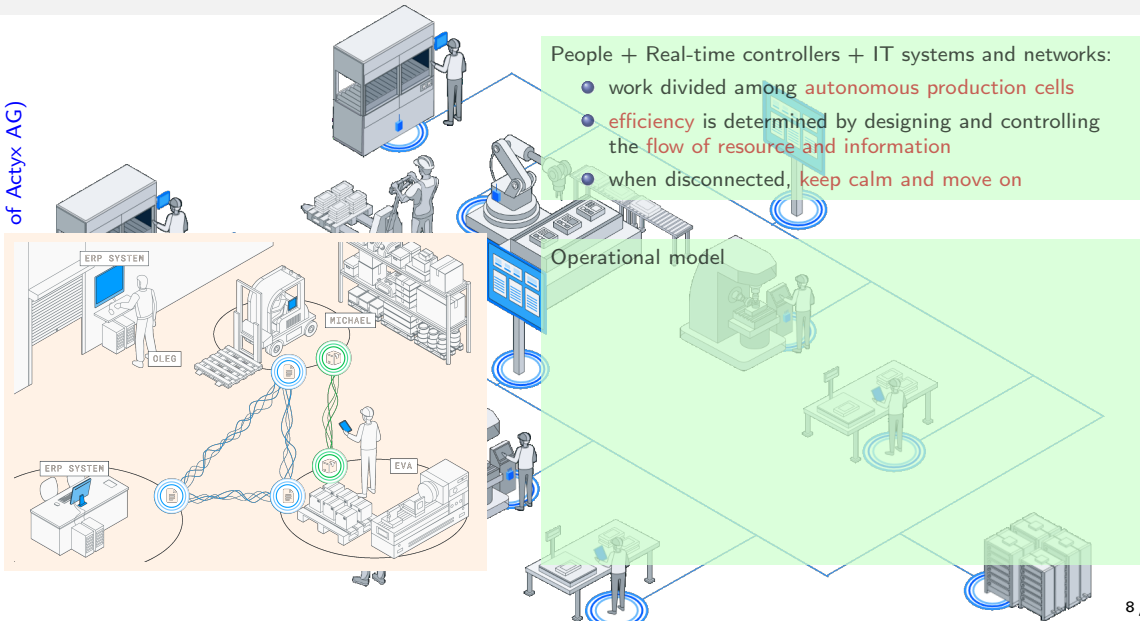
A collaborative environment and its execution model

(the pictures are courtesy of Actyx AG)



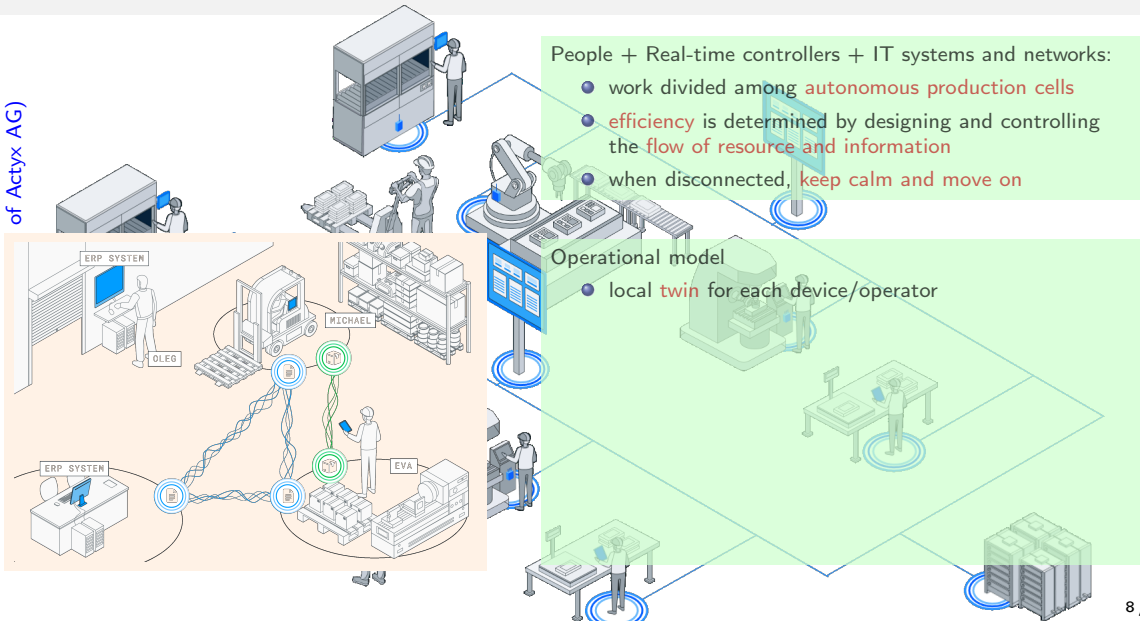
A collaborative environment and its execution model

of Actyx AG)



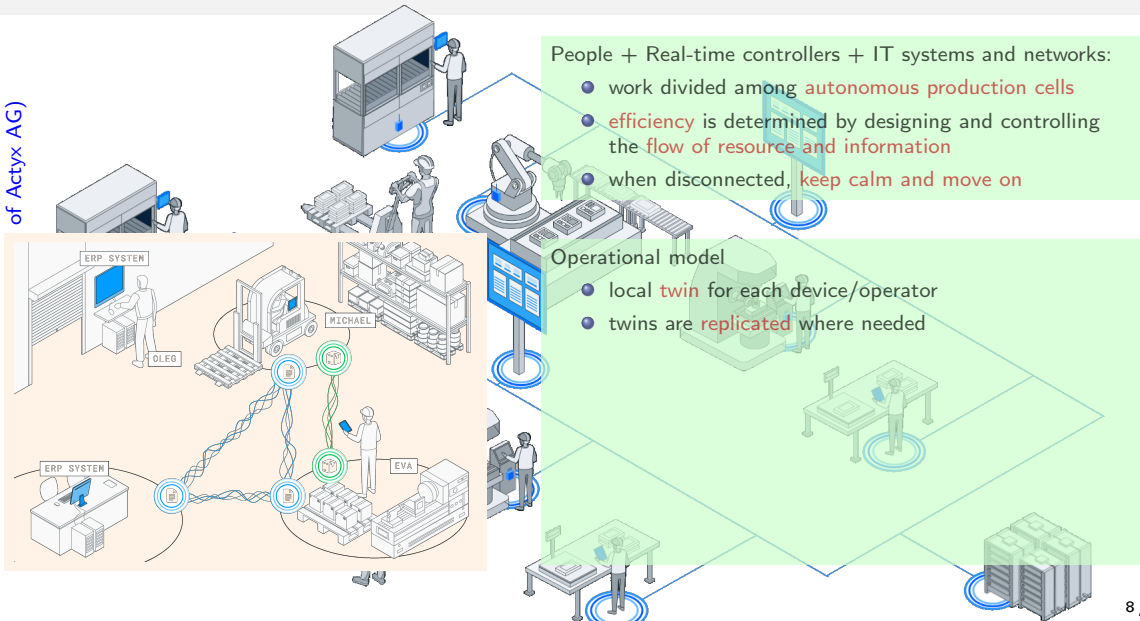
A collaborative environment and its execution model

of Actyx AG)



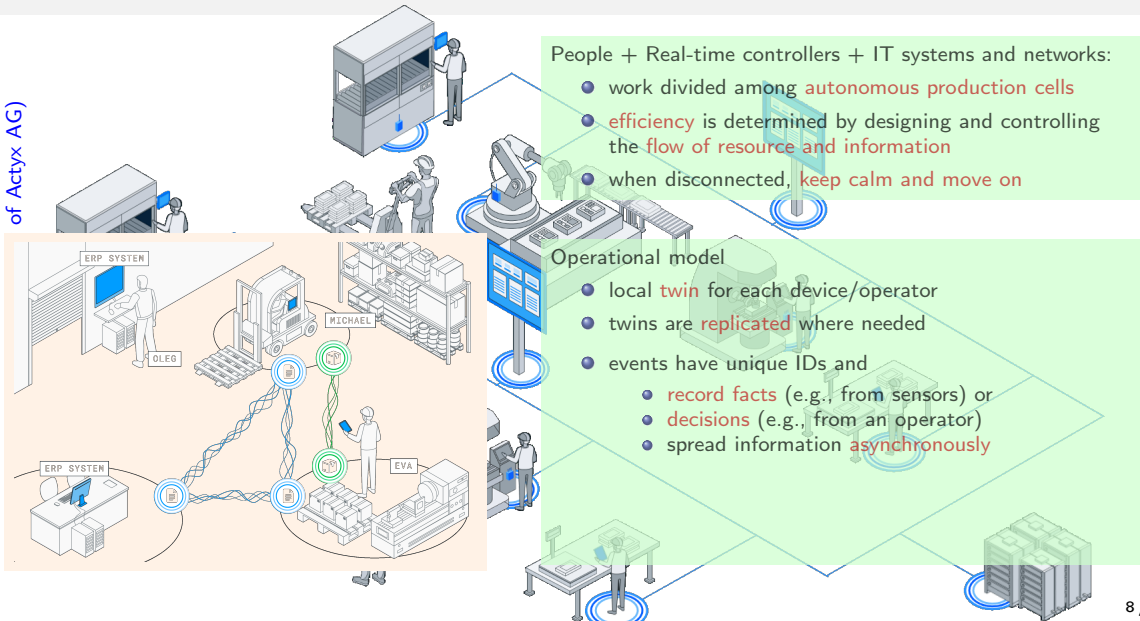
A collaborative environment and its execution model

of Actyx AG)



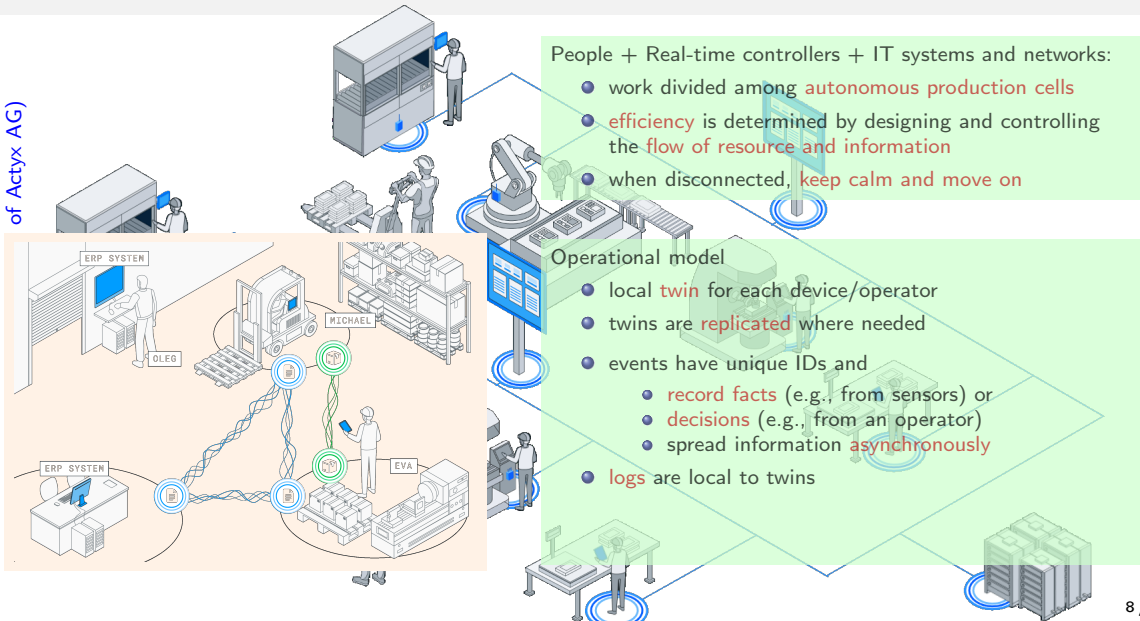
A collaborative environment and its execution model

of Actyx AG)



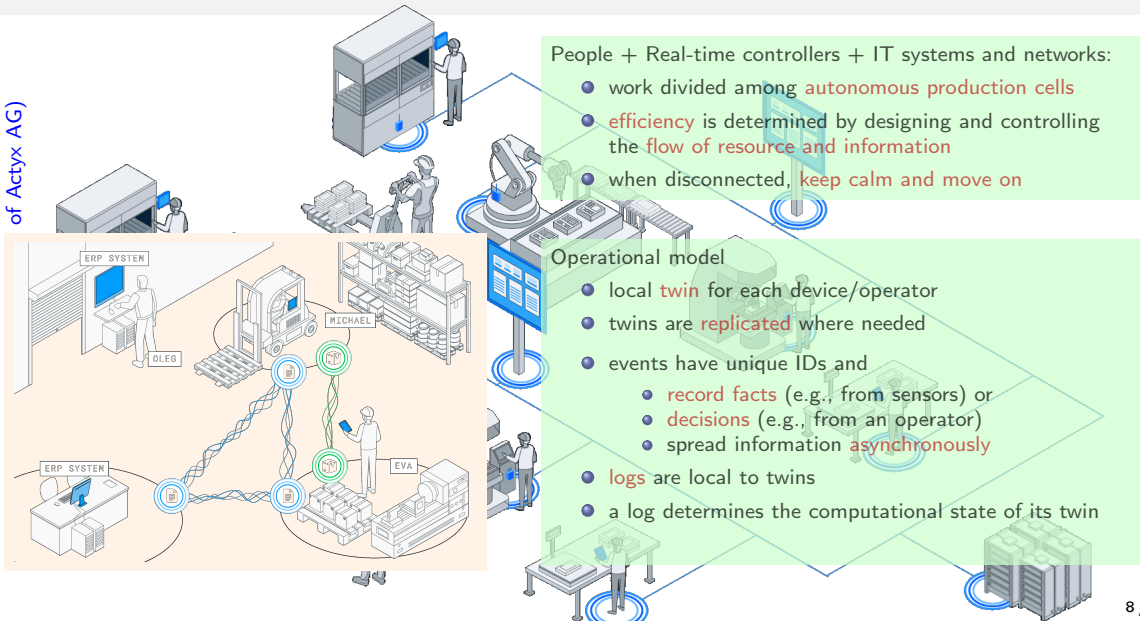
A collaborative environment and its execution model

of Actyx AG)



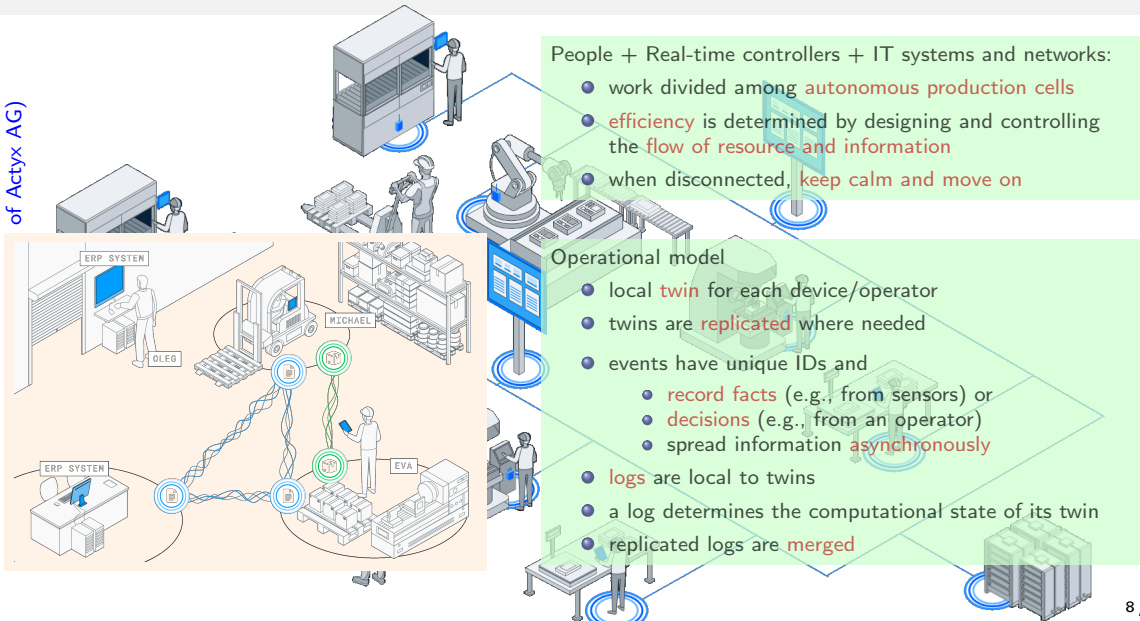
A collaborative environment and its execution model

of Actyx AG)



A collaborative environment and its execution model

of Actyx AG)



execute
+
propagate
+
merge

Other application domains / motivations

More applications

Robots (e.g., rescue missions or space applications)

Collaborative applications (<https://automerge.org/>)

Home automation

Other application domains / motivations

IoT...really?

Why your fridge and mobile **should go in the cloud** to talk to each other?

Other application domains / motivations

“Anytime, anywhere...” really?

like the AWS's outage on 25/11/2020

or almost all Google services down on 14/12/2020

DSL typical availability of 97% (& some SLA have no **lower bound**) checkout

<https://www.internetsociety.org/blog/2022/03/what-is-the-digital-divide>

Other application domains / motivations

Also, taking decisions locally

can reduce downtime

shifts data ownership

gets rid of any centralization point...for real

Plan of the talk

A motivating case study

Our formalisation

Our typing discipline

Tool support

Future work

– A formal model –

Events

e

Logs

$e_1 \cdot e_2 \dots$

Events

$\vdash e : t$

$src(e)$

Logs

$\vdash e_1 \cdot e_2 \dots : t_1 \cdot t_2 \dots$

Events $\vdash e : t$
 $src(e)$

Logs $\vdash e_1 \cdot e_2 \dots : t_1 \cdot t_2 \dots$

order induced by $\ell = e_1 \cdots e_n$ $e_i <_{\ell} e_j \iff i < j$

Ingredients (II): log shipping

Machine **Alice** **emits** logs upon **execution** of commands (we'll see how in a moment)

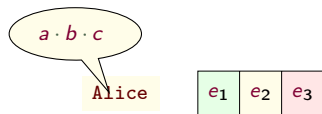
Ingredients (II): log shipping

Machine **Alice** **emits** logs upon **execution** of commands (we'll see how in a moment)
Such events are **appended** to the logs of machines in **two phases**:

Ingredients (II): log shipping

Machine **Alice** **emits** logs upon **execution** of commands (we'll see how in a moment)
Such events are **appended** to the logs of machines in **two phases**:

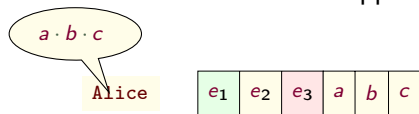
Phase I: emitted events are appended to the local log of the emitting machine



Ingredients (II): log shipping

Machine **Alice** **emits** logs upon **execution** of commands (we'll see how in a moment)
Such events are **appended** to the logs of machines in **two phases**:

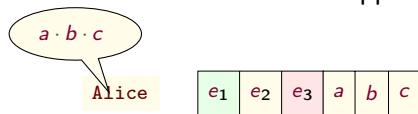
Phase I: emitted events are appended to the local log of the emitting machine



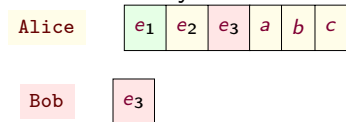
Ingredients (II): log shipping

Machine **Alice** **emits** logs upon **execution** of commands (we'll see how in a moment)
Such events are **appended** to the logs of machines in **two phases**:

Phase I: emitted events are appended to the local log of the emitting machine



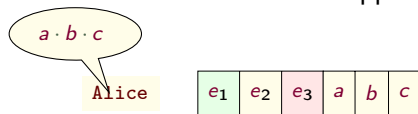
Phase II: newly emitted events are shipped to other machines



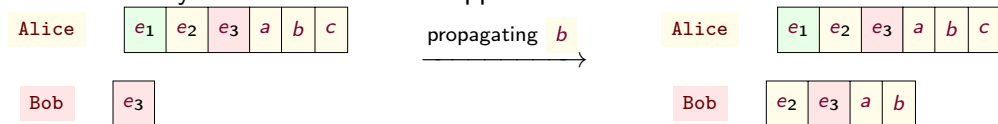
Ingredients (II): log shipping

Machine **Alice** **emits** logs upon **execution** of commands (we'll see how in a moment)
Such events are **appended** to the logs of machines in **two phases**:

Phase I: emitted events are appended to the local log of the emitting machine



Phase II: newly emitted events are shipped to other machines

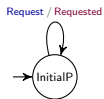


Machines by example




`InitialP` =

Machines by example



$\text{InitialP} = \text{Request} \mapsto \text{Requested}.$

Machines by example



$$\text{InitialP} = \text{Request} \mapsto \text{Requested} \cdot [\text{Requested?} \underline{\text{AuctionP}}]$$

Machines by example



$\text{InitialP} = \text{Request} \mapsto \text{Requested} \cdot [\text{Requested?} \underline{\text{AuctionP}}]$

$\text{AuctionP} =$

Machines by example



$\text{InitialP} = \text{Request} \mapsto \text{Requested} \cdot [\text{Requested?} \underline{\text{AuctionP}}]$

$\text{AuctionP} = [\text{Bid? BidderId?} \underline{\text{AuctionP}}]$

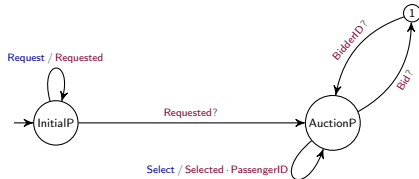
Machines by example



$\text{InitialP} = \text{Request} \mapsto \text{Requested} \cdot [\text{Requested?} \underline{\text{AuctionP}}]$

$\text{AuctionP} = [\text{Bid?} \underline{\text{BidderId? AuctionP}}]$

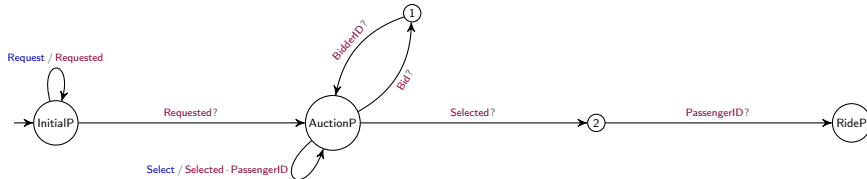
Machines by example



$\text{InitialP} = \text{Request} \mapsto \text{Requested} \cdot [\text{Requested?} \underline{\text{AuctionP}}]$

$\text{AuctionP} = \text{Select} \mapsto \text{Selected} \cdot \text{PassengerID} \cdot [\text{Bid? BidderID?} \underline{\text{AuctionP}}]$

Machines by example



InitialP = Request \mapsto Requested · [Requested? AuctionP]

AuctionP = Select \mapsto Selected · PassengerId · [
 Bid? BidderId? AuctionP
 &
 Selected? PassengerId? RideP
]

RideP = ...

Machines, formally

Fix a set of commands ranged over by c

Let κ range over finite maps from commands to non-empty log types

Machines, formally

Fix a set of commands ranged over by c

Let κ range over finite maps from commands to non-empty log types

Machine: **deterministic regular term** of $M \stackrel{\text{co}}{::=} \kappa \cdot [t_1?M_1 \& \dots \& t_n?M_n]$

Machines, formally

Fix a set of commands ranged over by c

Let κ range over finite maps from commands to non-empty log types

Machine: **deterministic regular term** of $M \stackrel{\text{co}}{::=} \kappa \cdot [t_1?M_1 \& \dots \& t_n?M_n]$

Think of machines as emitters/consumers of events with a semantics given in terms of state transition function :

$$\begin{aligned} \delta(M, \epsilon) &= M \\ \delta(M, e \cdot \ell) &= \begin{cases} \delta(M', \ell) & \text{if } \vdash e : t, M \xrightarrow{t?} M' \\ \delta(M, \ell) & \text{otherwise} \end{cases} \end{aligned}$$

That is

M with local log ℓ is in the implicit state $\delta(M, \ell)$ reached after processing each event in ℓ

Machines, formally

Fix a set of commands ranged over by c

Let κ range over finite maps from commands to non-empty log types

Machine: **deterministic regular term** of $M ::=^{\text{co}} \kappa \cdot [t_1?M_1 \& \dots \& t_n?M_n]$

Think of machines as emitters/consumers of events with a semantics given in terms of state transition function :

$$\delta(M, \epsilon) = M$$

$$\delta(M, e \cdot \ell) = \begin{cases} \delta(M', \ell) & \text{if } \vdash e : t, M \xrightarrow{t?} M' \\ \delta(M, \ell) & \text{otherwise} \end{cases}$$

$$\frac{\delta(M, \ell) \xrightarrow{c/1} \delta(M, \ell) \quad \ell' \text{ fresh} \quad \vdash \ell' : 1}{(M, \ell) \xrightarrow{c/1} (M, \ell \cdot \ell')}$$

That is

M with local log ℓ is in the implicit state $\delta(M, \ell)$ reached after processing each event in ℓ

That is

after processing the events in ℓ , M reaches a state enabling $c/1$ then the command execution can emit ℓ' of type 1 and append it to the local log of M

Swarms

Swarms: $M_1[\ell_1] \mid \dots \mid M_n[\ell_n] \mid \ell$ s.t. $\ell = \bigcup_{1 \leq i \leq n} \ell_i$ and $\ell_i \sqsubseteq \ell$ for $1 \leq i \leq n$

Swarms

Swarms: $M_1[\ell_1] \mid \dots \mid M_n[\ell_n] \mid \ell$ s.t. $\ell = \bigcup_{1 \leq i \leq n} \ell_i$ and $\ell_i \sqsubseteq \ell$ for $1 \leq i \leq n$

where $\ell_1 \sqsubseteq \ell_2$ is the sublog relation defined as

- $\ell_1 \subseteq \ell_2$ and $<_{\ell_1} \subseteq <_{\ell_2}$ and

- $e <_{\ell_2} e'$, $src(e) = src(e')$ and $e' \in \ell_1 \implies e \in \ell_1$

That is

all events of ℓ_1 appear in the same order in ℓ_2

That is

the per-source partitions of ℓ_1 are prefixes of the corresponding partitions of ℓ_2

Swarms

Swarms: $M_1[l_1] \mid \dots \mid M_n[l_n] \mid \ell$ s.t. $\ell = \bigcup_{1 \leq i \leq n} \ell_i$ and $\ell_i \sqsubseteq \ell$ for $1 \leq i \leq n$

where $\ell_1 \sqsubseteq \ell_2$ is the sublog relation defined as

- $\ell_1 \subseteq \ell_2$ and $<_{\ell_1} \subseteq <_{\ell_2}$ and

That is

all events of ℓ_1 appear in the same order in ℓ_2

- $e <_{\ell_2} e'$, $src(e) = src(e')$ and $e' \in \ell_1 \implies e \in \ell_1$

That is

the per-source partitions of ℓ_1 are prefixes of the corresponding partitions of ℓ_2

The propagation of newly generated events happens by merging logs:

Log merging: $\ell_1 \bowtie \ell_2 = \{\ell \mid \ell \subseteq \ell_1 \cup \ell_2 \text{ and } \ell_1 \sqsubseteq \ell \text{ and } \ell_2 \sqsubseteq \ell\}$

Semantics of swarms

By rule [Local] below, a command's execution updates both local and global logs

$$\frac{\begin{array}{l} S(i) = M[\ell_i] \quad M[\ell_i] \xrightarrow{c/1} M[\ell'_i] \quad \text{src}(\ell'_i \setminus \ell_i) = \{i\} \quad \ell' \in \ell \bowtie \ell'_i \end{array}}{(S, \ell) \xrightarrow{c/1} (S[i \mapsto M[\ell'_i]], \ell')} \text{[Local]}$$

Semantics of swarms

By rule [Local] below, a command's execution updates both local and global logs

$$\frac{S(i) = M[\ell_i] \quad M[\ell_i] \xrightarrow{c/1} M[\ell'_i] \quad \text{src}(\ell'_i \setminus \ell_i) = \{i\} \quad \ell' \in \ell \bowtie \ell'_i}{(S, \ell) \xrightarrow{c/1} (S[i \mapsto M[\ell'_i]], \ell')} \text{[Local]}$$

$$\frac{S(i) = M[\ell_i] \quad \ell_i \sqsubseteq \ell' \sqsubseteq \ell \quad \ell_i \subset \ell'}{(S, \ell) \xrightarrow{\tau} (S[i \mapsto M[\ell'_i]], \ell)} \text{[Prop]}$$

By rule [Prop] above, the propagation of events happens

- by shipping a **non-deterministically chosen** subset of events in the global log
- to a **non-deterministically chosen** machine

Plan of the talk

A motivating case study

Our formalisation

Our typing discipline

Tool support

Future work

– Behavioural types for swarms –

Inspired by choreographies

Quoting W3C:

*"[...] a **contract** [...] of the common **ordering conditions and constraints** under which **messages** are exchanged [...] from a **global viewpoint** [...]
Each **party** can then use the global definition to **build and test solutions** [...]
global specification is in turn **realised by combination of** the resulting **local systems**"*

Inspired by choreographies

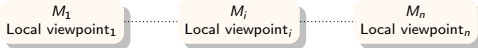
Quoting W3C:

*"[...] a **contract** [...] of the common **ordering conditions and constraints** under which **messages** are exchanged [...] from a **global viewpoint** [...]
Each **party** can then use the global definition to **build and test solutions** [...]
global specification is in turn **realised by combination of** the resulting **local systems**"*

Synchrony

Choreography G
global viewpoint

Asynchrony



Inspired by choreographies

Quoting W3C:

"[...] a *contract* [...] of the common *ordering conditions and constraints* under which *messages* are exchanged [...] from a *global viewpoint* [...]
Each party can then use the global definition to *build and test solutions* [...]
global specification is in turn *realised by combination of* the resulting *local systems*"

Synchrony

Choreography G
global viewpoint

Asynchrony

M_1
Local viewpoint₁

M_i
Local viewpoint_i

M_n
Local viewpoint_n

spec, no code

Inspired by choreographies

Quoting W3C:

"[...] a *contract* [...] of the common *ordering conditions and constraints* under which *messages* are exchanged [...] from a *global viewpoint* [...]
Each party can then use the global definition to *build and test solutions* [...]
global specification is in turn *realised by combination of* the resulting *local systems*"

Synchrony

Choreography G
global viewpoint

Well-formedness

Asynchrony

M_1
Local viewpoint₁

M_i
Local viewpoint_i

M_n
Local viewpoint_n


spec, no code

Inspired by choreographies

Quoting W3C:

"[...] a **contract** [...] of the common **ordering conditions and constraints** under which **messages** are exchanged [...] from a **global viewpoint** [...]"

Each party can then use the global definition to **build and test solutions** [...] global specification is in turn **realised by combination of** the resulting **local systems**"

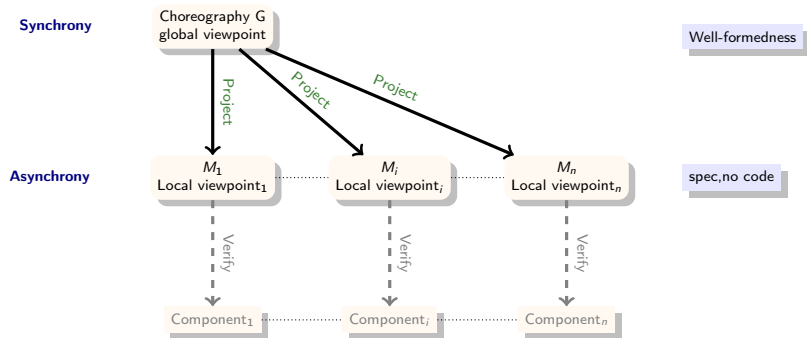


Inspired by choreographies

Quoting W3C:


"[...] a **contract** [...] of the common **ordering conditions and constraints** under which **messages** are exchanged [...] from a **global viewpoint** [...]"

Each party can then use the global definition to **build and test solutions** [...] global specification is in turn **realised by combination of** the resulting **local systems**"




Swarm protocols by example

An intuitive auction protocol for a passenger P to get a taxi T :



Swarm protocols by example

An intuitive auction protocol for a passenger P to get a taxi T :




We assume

- one passenger and one office (for simplicity)

Swarm protocols by example

An intuitive auction protocol for a passenger P to get a taxi T :




We assume

- one passenger and one office (for simplicity)
- but an arbitrary number of taxis

Swarm protocols by example

An intuitive auction protocol for a passenger P to get a taxi T :



We assume

- one passenger and one office (for simplicity)
- but an arbitrary number of taxis
- a receipt is issued by the office O at the end of the ride (if any)

Swarm protocols: global type for local-first applications

An **idealised** specification relying on **synchronous communication**

Fix a set of roles ranged over by \mathbf{R} (e.g., \mathbf{P} , \mathbf{T} , and \mathbf{O} on slide 31)

The syntax of swarm protocols is again given co-inductively:

$$\mathbf{G} ::=^{\text{co}} \sum_{i \in I} \mathbf{c}_i @ \mathbf{R}_i \langle \mathbf{l}_i \rangle . \mathbf{G}_i \quad | \quad 0 \quad \text{where } I \text{ is a finite set (of indexes)}$$

An example

A swarm protocol for the taxi scenario on slide 31:

$$G = \text{Request@P}\langle \text{Requested} \rangle . G_{\text{auction}}$$

$$\begin{aligned} G_{\text{auction}} &= \text{Offer@T}\langle \text{Bid} \cdot \text{BidderID} \rangle . G_{\text{auction}} \\ &+ \text{Select@P}\langle \text{Selected} \cdot \text{PassengerID} \rangle . G_{\text{choose}} \end{aligned}$$

$$\begin{aligned} G_{\text{choose}} &= \text{Arrive@T}\langle \text{Arrived} \rangle . \text{Start@P}\langle \text{Started} \rangle . G_{\text{ride}} \\ &+ \text{Cancel@P}\langle \text{Cancelled} \rangle . \text{Receipt@O}\langle \text{Receipt} \rangle . 0 \end{aligned}$$

$$\begin{aligned} G_{\text{ride}} &= \text{Record@T}\langle \text{Path} \rangle . G_{\text{ride}} \\ &+ \text{Finish@P}\langle \text{Finished} \cdot \text{Rating} \rangle . \text{Receipt@O}\langle \text{Receipt} \rangle . 0 \end{aligned}$$

An example

A swarm protocol for the taxi scenario on slide 31:

$$G = \text{Request@P}\langle \text{Requested} \rangle . G_{\text{auction}}$$

$$\begin{aligned} G_{\text{auction}} &= \text{Offer@T}\langle \text{Bid} \cdot \text{BidderID} \rangle . G_{\text{auction}} \\ &+ \text{Select@P}\langle \text{Selected} \cdot \text{PassengerID} \rangle . G_{\text{choose}} \end{aligned}$$

*Note the log types
in each prefixes*

$$\begin{aligned} G_{\text{choose}} &= \text{Arrive@T}\langle \text{Arrived} \rangle . \text{Start@P}\langle \text{Started} \rangle . G_{\text{ride}} \\ &+ \text{Cancel@P}\langle \text{Cancelled} \rangle . \text{Receipt@O}\langle \text{Receipt} \rangle . 0 \end{aligned}$$

$$\begin{aligned} G_{\text{ride}} &= \text{Record@T}\langle \text{Path} \rangle . G_{\text{ride}} \\ &+ \text{Finish@P}\langle \text{Finished} \cdot \text{Rating} \rangle . \text{Receipt@O}\langle \text{Receipt} \rangle . 0 \end{aligned}$$

Swarm protocols as FSA

Like for machines, a swarm protocols $G = \sum_{i \in I} c_i @ R_i \langle 1_i \rangle$. G_i has an associated FSA:

- the set of states consists of G plus the states in G_i for each $i \in \{1 \dots, n\}$
- G is the initial state
- for each $i \in I$, G has a transition to state G_i labelled with $c_i @ R_i \langle 1_i \rangle$, written
$$G \xrightarrow{c_i / 1_i} G_i$$

Semantics of swarm protocols

One rule only!

$$\frac{}{(G, \ell) \xrightarrow{c/1} (G, \ell)} \text{[G-Cmd]}$$

Semantics of swarm protocols

One rule only!

$$\frac{\delta(\mathbf{G}, \ell) \xrightarrow{\mathbf{c}/\mathbf{l}} \mathbf{G}'}{(\mathbf{G}, \ell) \xrightarrow{\mathbf{c}/\mathbf{l}} (\mathbf{G}, \ell)} \text{ [G-Cmd]}$$

where

$$\delta(\mathbf{G}, \ell) = \begin{cases} \mathbf{G} & \text{if } \ell = \epsilon \\ \delta(\mathbf{G}', \ell'') & \text{if } \mathbf{G} \xrightarrow{\mathbf{c}/\mathbf{l}} \mathbf{G}' \text{ and } \vdash \ell' : \mathbf{l} \text{ and } \ell = \ell' \cdot \ell'' \\ \perp & \text{otherwise} \end{cases}$$

*Logs to be consumed "atomically",
hence $\delta(\mathbf{G}, \ell)$ may be undefined*

Semantics of swarm protocols

One rule only!

$$\frac{\delta(\mathbf{G}, \ell) \xrightarrow{c/1} \mathbf{G}' \quad \vdash \ell' : 1 \quad \ell' \text{ log of fresh events}}{(\mathbf{G}, \ell) \xrightarrow{c/1} (\mathbf{G}, \ell \cdot \ell')} \text{[G-Cmd]}$$

where

$$\delta(\mathbf{G}, \ell) = \begin{cases} \mathbf{G} & \text{if } \ell = \epsilon \\ \delta(\mathbf{G}', \ell'') & \text{if } \mathbf{G} \xrightarrow{c/1} \mathbf{G}' \text{ and } \vdash \ell' : 1 \text{ and } \ell = \ell' \cdot \ell'' \\ \perp & \text{otherwise} \end{cases}$$

*Logs to be consumed "atomically",
hence $\delta(\mathbf{G}, \ell)$ may be undefined*

Semantics of swarm protocols

One rule only!

$$\frac{\delta(\mathbf{G}, \ell) \xrightarrow{\text{c}/\mathbf{1}} \mathbf{G}' \quad \vdash \ell' : \mathbf{1} \quad \ell' \text{ log of fresh events}}{(\mathbf{G}, \ell) \xrightarrow{\text{c}/\mathbf{1}} (\mathbf{G}, \ell \cdot \ell')} \text{[G-Cmd]}$$

where

$$\delta(\mathbf{G}, \ell) = \begin{cases} \mathbf{G} & \text{if } \ell = \epsilon \\ \delta(\mathbf{G}', \ell'') & \text{if } \mathbf{G} \xrightarrow{\text{c}/\mathbf{1}} \mathbf{G}' \text{ and } \vdash \ell' : \mathbf{1} \text{ and } \ell = \ell' \cdot \ell'' \\ \perp & \text{otherwise} \end{cases}$$

Logs to be consumed "atomically", hence $\delta(\mathbf{G}, \ell)$ may be undefined

We restrict ourselves to deterministic swarm protocols that is, on different transitions from a same state

- log types start differently
- pairs (command,role) differ

log determinism

command determinism

From swarm protocols to machines

Transitions of a swarm protocol G are labelled with a role that may invoke the command

From swarm protocols to machines

Transitions of a swarm protocol G are labelled with a role that may invoke the command

Each machine plays one role

From swarm protocols to machines

Transitions of a swarm protocol G are labelled with a role that may invoke the command

Each machine plays one role



Obtain machines by projecting G on each role

From swarm protocols to machines

Transitions of a swarm protocol G are labelled with a role that may invoke the command

Each machine plays one role



Obtain machines by projecting G on each role

First attempt

$$\left(\sum_{i \in I} c_i @ R_i \langle l_i \rangle \cdot G_i \right) \downarrow_R = \kappa \cdot [\&_{i \in I} l_i ? G_i \downarrow_R]$$

where $\kappa = \{(c_i / l_i) \mid R_i = R \text{ and } i \in I\}$

From swarm protocols to machines

Transitions of a swarm protocol G are labelled with a role that may invoke the command

Each machine plays one role



Obtain machines by projecting G on each role

First attempt

$$\left(\sum_{i \in I} c_i @ R_i \langle l_i \rangle \cdot G_i \right) \downarrow_R = \kappa \cdot [\&_{i \in I} l_i? G_i \downarrow_R]$$

where $\kappa = \{(c_i / l_i) \mid R_i = R \text{ and } i \in I\}$

simple, but

- projected machines are large in all but the most trivial cases
- processing **all** events is undesirable: security and efficiency

Another attempt



Let's subscribe to subscriptions : maps from roles to sets of event types

*In pub-sub,
processes subscribe
to "topics"*

Another attempt



Let's subscribe to subscriptions : maps from roles to sets of event types

*In pub-sub,
processes subscribe
to "topics"*

Given $G = \sum_{i \in I} c_i @ R_i \langle 1_i \rangle . G_i$, the
projection of G on a role R with respect to subscription σ is

$$G \downarrow_R^\sigma = \kappa \cdot [\&_{j \in J} \text{filter}(1_j, \sigma(R)) ? G_j \downarrow_R^\sigma]$$

where

Another attempt



Let's subscribe to subscriptions : maps from roles to sets of event types

*In pub-sub,
processes subscribe
to "topics"*

Given $G = \sum_{i \in I} c_i @ R_i \langle 1_i \rangle . G_i$, the
projection of G on a role R with respect to subscription σ is

$$G \downarrow_R^\sigma = \kappa \cdot [\&_{j \in J} \text{filter}(1_j, \sigma(R)) ? G_j \downarrow_R^\sigma] \quad \text{where}$$

$$\kappa = \{c_i / 1_i \mid R_i = R \text{ and } i \in I\}$$

$$J = \{i \in I \mid \text{filter}(1_i, \sigma(R)) \neq \epsilon\}$$

$$\text{filter}(1, E) = \begin{cases} \epsilon, & \text{if } t = \epsilon \\ t \cdot \text{filter}(1', E) & \text{if } t \in E \text{ and } 1 = t \cdot 1' \\ \text{filter}(1, E) & \text{otherwise} \end{cases}$$

Well-formedness

Trading consistency for availability has implications:

Well-formedness = Causality

Trading consistency for availability has implications:

Propagation of events is non-atomic (cf. rule [Prop])

\Rightarrow differences in how machines perceive the (state of the) computation

Causality

Fix a subscription σ . For each branch $i \in I$ of $G = \sum_{i \in I} c_i @ R_i \langle 1_i \rangle . G_i$

Explicit re-enabling $\sigma(R_i) \cap 1_i \neq \emptyset$

If R should have c enabled after c' then $\sigma(R)$ contains some event type emitted by c'

Command causality if R executes a command in G_i
then $\sigma(R) \cap 1_i \neq \emptyset$ and $\sigma(R) \cap 1_i \supseteq \bigcup_{R' \in \sigma G_i} \sigma(R') \cap 1_i$

Well-formedness = Causality + Determinacy

Trading consistency for availability has implications:

Propagation of events is non-atomic (cf. rule [Prop])

\implies different roles may take inconsistent decisions

Causality & Determinacy

Fix a subscription σ . For each branch $i \in I$ of $G = \sum_{i \in I} c_i @ R_i \langle 1_i \rangle . G_i$

Explicit re-enabling $\sigma(R_i) \cap 1_i \neq \emptyset$

Command causality if R executes a command in G_i
then $\sigma(R) \cap 1_i \neq \emptyset$ and $\sigma(R) \cap 1_i \supseteq \bigcup_{R' \in_\sigma G_i} \sigma(R') \cap 1_i$

Determinacy $R \in_\sigma G_i \implies 1_i[0] \in \sigma(R)$

Well-formedness = Causality + Determinacy - Confusion

Trading consistency for availability has implications:

Propagation of events is non-atomic (cf. rule [Prop])

\Rightarrow branches unambiguously identified and events emitted on eventually discharged branches ignored

Causality & Determinacy & Confusion freeness

Fix a subscription σ . For each branch $i \in I$ of $G = \sum_{i \in I} c_i @ R_i \langle 1_i \rangle . G_i$

Explicit re-enabling $\sigma(R_i) \cap 1_i \neq \emptyset$

Command causality if R executes a command in G_i
then $\sigma(R) \cap 1_i \neq \emptyset$ and $\sigma(R) \cap 1_i \supseteq \bigcup_{R' \in_\sigma G_i} \sigma(R') \cap 1_i$

Determinacy $R \in_\sigma G_i \Rightarrow 1_i[0] \in \sigma(R)$

Confusion freeness there is a unique subtree G' of G emitting t
for each t starting a log emitted by a command in G

Implementations

Write $\ell \equiv_{\mathbf{G}, \sigma} \ell'$ when ℓ and ℓ' have the same effective type wrt \mathbf{G} and σ .

A swarm (\mathbf{S}, ϵ) is eventually faithful to \mathbf{G} and σ if $(\mathbf{S}, \epsilon) \Longrightarrow (\mathbf{S}, \ell)$ then there is $(\mathbf{G}, \epsilon) \Longrightarrow (\mathbf{G}, \ell')$ with $\ell \equiv_{\mathbf{G}, \sigma} \ell'$

Implementations

Write $\ell \equiv_{\mathbf{G}, \sigma} \ell'$ when ℓ and ℓ' have the same effective type wrt \mathbf{G} and σ .

A swarm (\mathbf{S}, ϵ) is eventually faithful to \mathbf{G} and σ if $(\mathbf{S}, \epsilon) \Longrightarrow (\mathbf{S}, \ell)$ then there is $(\mathbf{G}, \epsilon) \Longrightarrow (\mathbf{G}, \ell')$ with $\ell \equiv_{\mathbf{G}, \sigma} \ell'$

A (σ, \mathbf{G}) -realisation is a swarm (\mathbf{S}, ϵ) of size n such that, for each $1 \leq i \leq n$, there exists a role $\mathbf{R} \in \text{roles}(\mathbf{G}, \sigma)$ such that $\mathbf{S}(i) = \mathbf{G} \downarrow_{\mathbf{R}}^{\sigma} []$

Implementations & projections

Write $\ell \equiv_{G,\sigma} \ell'$ when ℓ and ℓ' have the same effective type wrt G and σ .


A swarm (S, ϵ) is eventually faithful to G and σ if $(S, \epsilon) \Longrightarrow (S, \ell)$ then there is $(G, \epsilon) \Longrightarrow (G, \ell')$ with $\ell \equiv_{G,\sigma} \ell'$

A (σ, G) -realisation is a swarm (S, ϵ) of size n such that, for each $1 \leq i \leq n$, there exists a role $R \in \text{roles}(G, \sigma)$ such that $S(i) = G \downarrow_R^\sigma$

Lemma (Projections of well-formed protocols are eventually faithful)


If G is a σ -WF protocol and $(\delta(G \downarrow_R^\sigma, \ell)) \downarrow_{c/1}$ then there exists $\ell' \equiv_{G,\sigma} \ell$ such that $(G, \epsilon) \Longrightarrow (G, \ell')$ and $\delta(G, \ell') \xrightarrow{c/1} G'$

On correct realisations



*A set of runs is consistent
when its elements are
pair-wise consistent*

On correct realisations




A set of runs is consistent when its elements are pair-wise consistent

Notation

For $(G, \epsilon) \xrightarrow{c_1 / l_1} (G, \ell_1) \xrightarrow{c_2 / l_2} \dots \xrightarrow{c_n / l_n} (G, \overbrace{\ell_1 \cdot \ell_2 \cdot \dots \cdot \ell_n}^{= \ell})$
 let $\ell^{(j)} = \ell_j \cdot \dots \cdot \ell_1$

On correct realisations



A set of runs is consistent when its elements are pair-wise consistent

Notation

For $(G, \epsilon) \xrightarrow{c_1 / l_1} (G, \ell_1) \xrightarrow{c_2 / l_2} \dots \xrightarrow{c_n / l_n} (G, \overbrace{\ell_1 \cdot \ell_2 \cdot \dots \cdot \ell_n}^{=\ell})$
 let $\ell^{(j)} = \ell_j \cdot \dots \cdot \ell_1$

Admissibility

A log ℓ is admissible for a σ -WF protocol G if there are consistent runs $\{(G, \epsilon) \implies (G, \ell_i)\}_{1 \leq i \leq k}$ and a log $\ell' \in (\bowtie_{1 \leq i \leq k} \ell_i)$ such that $\ell = \bigcup_{1 \leq i \leq k} \ell_i$ and

$$\ell' \equiv_{G, \sigma} \ell \quad \text{and} \quad \ell_i^{(j)} \sqsubseteq \ell \text{ for all } 1 \leq i \leq k$$

Hereafter, G denotes a σ -WF protocol

Results

Lemma (Well-formedness generates any admissible log)

If ℓ is admissible for G then there is a log ℓ' such that $(G, \epsilon) \Longrightarrow (G, \ell')$ and $\ell \equiv_{G, \sigma} \ell'$

Lemma (Admissibility is preserved)

Let ℓ_1 and $\ell_2 \subseteq \ell_1$ be admissible logs for G . If $(G, \ell_2) \xrightarrow{c/1} (G, \ell_2 \cdot \ell_3)$ and $\ell \in \ell_1 \bowtie (\ell_2 \cdot \ell_3)$ then ℓ is admissible for G

Theorem (Well-formed protocols generate only admissible logs)

If $(S, \epsilon) \Longrightarrow (S', \ell)$ for (S, ϵ) realisation of G then ℓ is admissible for G

Corollary

Every realisation of G is eventually faithful wrt G and σ

On complete realisations

Complete realisations

A (σ, G) -realisation (S, ϵ) of size n is complete if for all $R \in \text{roles}(G, \sigma)$ there exists $1 \leq i \leq n$ such that $S(i) = G \downarrow_R^\sigma$

Lemma (Projections reflect swarm protocols)

If $(G, \epsilon) \implies (G, \ell)$ then $\delta(G \downarrow_R^\sigma, \ell) = \delta(G, \ell) \downarrow_R^\sigma$ for all $R \in \text{roles}(G, \sigma)$

Theorem (Complete realisations reflect the protocol)

Let (S, ϵ) be a complete realisation of G . If $(G, \epsilon) \implies (G, \ell)$ then there is a swarm S' such that $(S, \epsilon) \implies (S', \ell)$

Plan of the talk

A motivating case study

Our formalisation

Our typing discipline

Tool support

Future work

– Tooling –

```
// analogous for other events; "type" property matches type name (checked by tool)
type Requested = { type: 'Requested'; pickup: string; dest: string }
type Events = Requested | Bid | BidderID | Selected | ...
```

```
/** Initial state for role P */
```

```
@proto('taxiRide') // decorator injects inferred protocol into runtime
```


```
export class InitialP extends State<Events> {
  constructor(public id: string) { super() }
  execRequest(pickup: string, dest: string) {
    return this.events({ type: 'Requested', pickup, dest })
  }
  onRequest(ev: Requested) {
    return new AuctionP(this.id, ev.pickup, ev.dest, [])
  }
}
```

```
@proto('taxiRide')
```

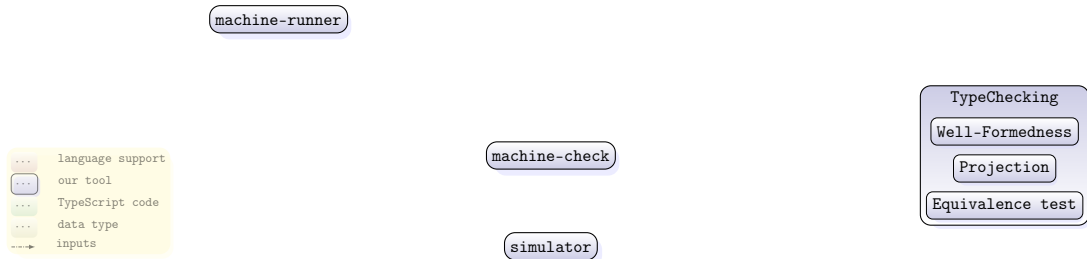
```
export class AuctionP extends State<Events> {
  constructor(public id: string, public pickup: string, public dest: string,
    public bids: BidData[]) { super() }
  onBid(ev1: Bid, ev2: BidderID) {
    const [ price, time ] = ev1
    this.bids.push({ price, time, bidderID: ev2.id })
    return this
  }
  execSelect(taxiId: string) {
    return this.events({ type: 'Selected', taxiID },
      { type: 'PassengerID', id: this.id })
  }
  onSelected(ev: Selected, id: PassengerID) {
    return new RideP(this.id, ev.taxiID)
  }
}
```

```
@proto('taxiRide')
```

```
export class RideP extends State<Events> { ... }
```




Architecture




- TypeChecking implements the functionalities of our typing discipline
- simulator simulates the semantics of swarm realisations
- machine-check and machine-runner integrate our framework in the Actyx platform

Architecture



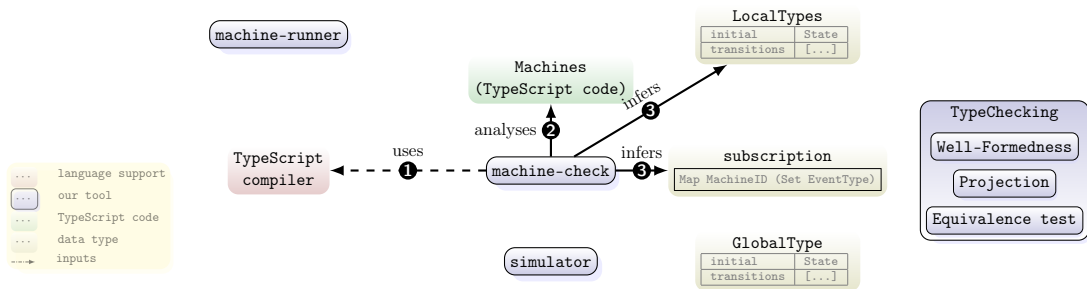
- TypeChecking implements the functionalities of our typing discipline
- simulator simulates the semantics of swarm realisations
- machine-check and machine-runner integrate our framework in the Actyx platform

Architecture




- TypeChecking implements the functionalities of our typing discipline
- simulator simulates the semantics of swarm realisations
- machine-check and machine-runner integrate our framework in the Actyx platform

Architecture



- TypeChecking implements the functionalities of our typing discipline
- simulator simulates the semantics of swarm realisations
- machine-check and machine-runner integrate our framework in the Actyx platform

Architecture



- TypeChecking implements the functionalities of our typing discipline
- simulator simulates the semantics of swarm realisations
- machine-check and machine-runner integrate our framework in the Actyx platform

If you want to play with our prototype?

Have a look at

- our ECOOP artifact paper (not online yet; extended version at <https://arxiv.org/abs/2305.04848>)
- code at <https://doi.org/10.5281/zenodo.7737188>
- An ISSTA tool paper from Actyx (<https://arxiv.org/abs/2306.09068>)

Plan of the talk

A motivating case study

Our formalisation

Our typing discipline

Tool support

Future work

– Epilogue –

To be continued....

There are a number of future directions to explore:

To be continued....

There are a number of future directions to explore:

Identify weaker conditions for well-formedness

To be continued....

There are a number of future directions to explore:

Identify weaker conditions for well-formedness

“Efficiency”

To be continued....

There are a number of future directions to explore:

- Identify weaker conditions for well-formedness

- “Efficiency”

- Subscriptions are hard to determine

To be continued....

There are a number of future directions to explore:

- Identify weaker conditions for well-formedness

- “Efficiency”

- Subscriptions are hard to determine

- Relax some of our assumptions

To be continued....

There are a number of future directions to explore:

- Identify weaker conditions for well-formedness

- “Efficiency”

- Subscriptions are hard to determine

- Relax some of our assumptions

 - Compensations

 - Unreliable propagation

To be continued....

There are a number of future directions to explore:

- Identify weaker conditions for well-formedness

- “Efficiency”

- Subscriptions are hard to determine

- Relax some of our assumptions

 - Compensations

 - Unreliable propagation

 - Adversarial contexts

To be continued....

There are a number of future directions to explore:

- Identify weaker conditions for well-formedness

- “Efficiency”

- Subscriptions are hard to determine

- Relax some of our assumptions

 - Compensations

 - Unreliable propagation

 - Adversarial contexts

-

Summary

An interesting paradigm grounded on principles for local-first principles: temporary inconsistency are tolerated provided that they can be (and are) resolved at some point

Summary

An interesting paradigm grounded on principles for local-first principles: temporary inconsistency are tolerated provided that they can be (and are) resolved at some point

A formal semantics that faithfully captures Actyx's platform

An interesting paradigm grounded on principles for local-first principles: temporary inconsistency are tolerated provided that they can be (and are) resolved at some point

A formal semantics that faithfully captures Actyx's platform

and behavioural types to specify and verify eventual consensus

Thank you!